

AI & Engineering — Strategy on a Page

Sandbox architecture · GitHub Copilot rollout · Claude Code on Foundry parallel pilot · Governance

Owner: Head of AI & Automation
Stack: Microsoft + AI Foundry
Dev org: <200 engineers · Africa · Fintech
Version: 1.0 · Apr 2026

MISSION Unlock measurable engineering productivity and safe AI adoption under a single, regulator-defensible control plane — on the stack we already own.

Sandbox & Dev Environment

PILLAR 01

Isolated Foundry + Dev Box + Azure Deployment Environments, preloaded with Purview DLP, CMK and synthetic data.

< 1 day

ENV PROVISIONING

100%

CMK + DLP ON

GitHub Copilot Enterprise

PILLAR 02

8-12 week pilot, 40-60 engineers. Inline completion, Chat, Code Review, Autofix. Scale to full org on evidence.

+15-30%

THROUGHPUT

>25%

ACCEPT RATE

Claude Code on Foundry

PILLAR 03

Parallel 10-15 seat sub-cohort. Same metrics, head-to-head. MACC-billed. Reduces single-vendor concentration.

Same


BASELINE METRICS

MACC

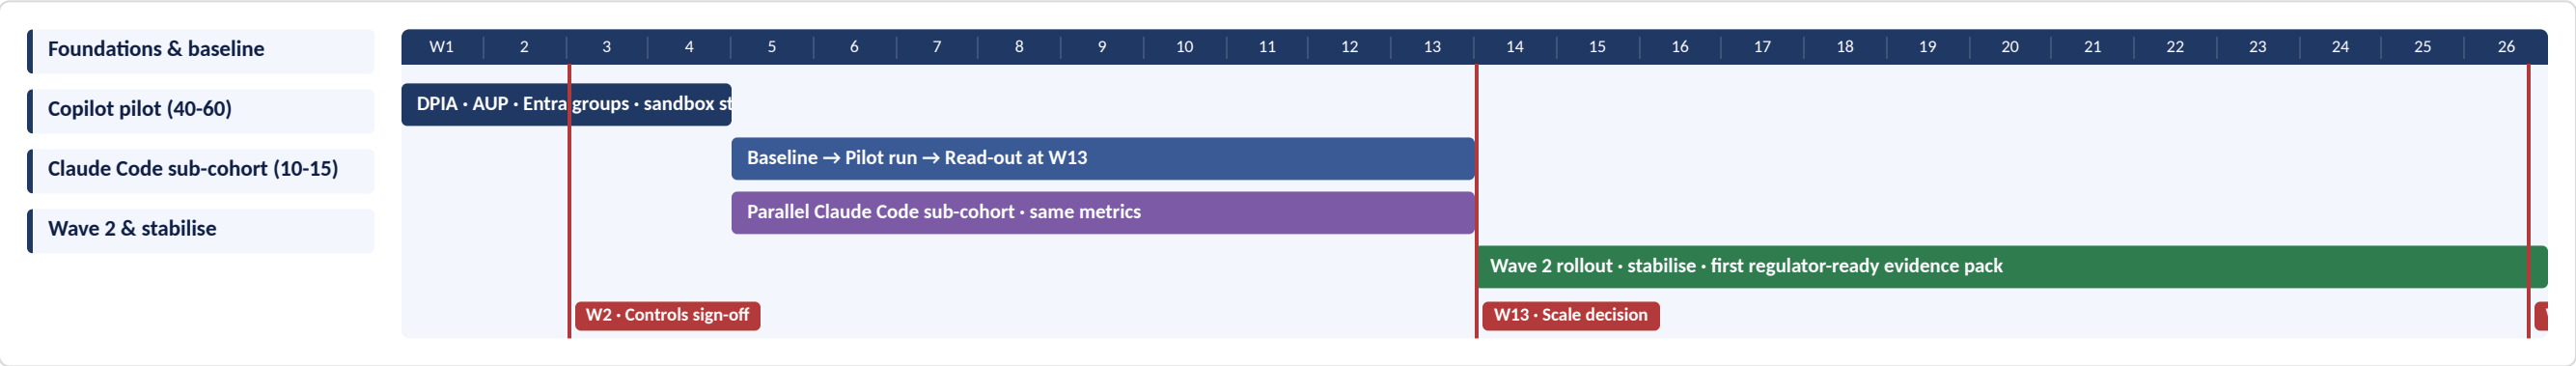
AZURE BILLING

ARCHITECTURE · CONTROL PLANE FROM DEVELOPER TO PRODUCTION

Layer	Services
Developer surface	VS Code · JetBrains · Dev Box · Codespaces
AI coding tools	GitHub Copilot Enterprise · Claude Code (Foundry)
Models & orchestration	Azure OpenAI · Claude Sonnet 4.6 / Opus 4.7 · Foundry Agents
Code & CI/CD	GitHub Enterprise Cloud · GHAS · Azure DevOps pipelines
Identity & network	Entra ID · Conditional Access · Private Link · CMK
Data & DLP	Purview classification & policies · synthetic PII
Monitor & respond	Defender for Cloud · Sentinel · Azure Monitor
Governance	AI Council · Use-case register · ISO 42001 / NIST AI RMF



ROLLOUT · 26-WEEK PLAN WITH DECISION GATES



Week-13 decision tree · choosing the agentic-coding tool strategy

Week-13 read-out

Copilot vs Claude Code · same metrics · same cohort

Copilot-primary

OPTION A

Better cost-per-accepted-change and engineer preference. Claude retained for narrow expert cohort on long-horizon refactors.

Claude-primary

OPTION B

Wins on task completion, time-to-done, preference. Copilot retained for inline completion & Code Review flows.

Dual-use

MOST LIKELY

Copilot for inline + Code Review + Chat. Claude Code for agentic multi-file changes, migrations, deep refactors.

Stop one track

SAFETY

One tool fails safety, IP or adoption thresholds. De-scope; the other proceeds to Wave 2 alone.

Guardrails in every path

- Prompts & outputs not used to train foundation models (contractual)
- Public-code match filter ON · IP indemnification
- Content exclusions on payments, fraud, secrets repos
- Agentic execution confined to Dev Box / Codespaces
- No production credentials in any agent context

Approvals before seat #1

- DPIA signed by DPO (POPIA / NDPA / KDPA as applicable)
- CISO sign-off on monitoring & agentic threat model
- Legal: AUP & contract addendum
- CRO: model-risk registration
- Procurement: MACC drawdown & cost alerts

What the board sees quarterly

- DORA: lead time · deploy freq · change-fail · MTTR
- Adoption: seats active · accept rate · cohort coverage
- Value: time saved · cost per accepted change
- Safety: secret-scan regressions · CodeQL trend · incidents
- Governance: use-cases in register · audit findings closed